

ARM Cortex-M Basics:

Core Components of Cortex-M4

1. **Processor Core**
 - **3-stage pipeline Harvard architecture** for efficient instruction execution.
 - Supports **Thumb-2 instruction set**, combining high code density with 32-bit performance.
 - Includes a **Floating-Point Unit (FPU)** (optional) for IEEE754-compliant single-precision floating-point operations.
2. **Memory Protection Unit (MPU)** (Optional)
 - Provides up to 8 memory regions for defining access permissions and attributes.
 - Enhances system reliability by preventing unauthorized access to critical memory areas.
3. **Nested Vectored Interrupt Controller (NVIC)**
 - Low-latency interrupt handling with up to **256 priority levels**.
 - Supports tail-chaining for efficient interrupt processing.
 - Integrated with sleep modes for power management.
4. **System Control Block (SCB)**
 - Manages system exceptions, configuration, and control.
 - Includes registers for fault handling, sleep modes, and system control.
5. **System Timer (SysTick)**
 - A 24-bit count-down timer, commonly used as an RTOS tick timer.
6. **Debug and Trace Features**
 - **Serial Wire Debug (SWD)** and **Serial Wire Viewer (SWV)** for reduced pin-count debugging.
 - Optional **Embedded Trace Macrocell (ETM)** for full instruction trace.
 - **Flash Patch and Breakpoint Unit (FPB)** for hardware breakpoints and code patching.
7. **Bus Interfaces**
 - **AMBA (Advanced Microcontroller Bus Architecture)** for high-speed, low-latency memory access.
 - Supports **unaligned data access** and **atomic bit manipulation**.

Important Features of Cortex-M4

1. **High Performance**
 - Efficient 3-stage pipeline with **single-cycle multiply and accumulate (MAC)** operations.
 - **Hardware divide (SDIV/UDIV)** for fast arithmetic operations.
2. **Low Power Consumption**
 - Integrated sleep modes (**Sleep and Deep Sleep**) for power optimization.
 - Wake-up mechanisms via interrupts or events.
3. **Digital Signal Processing (DSP) Capabilities**
 - **Saturating arithmetic** for signal processing.

- **SIMD (Single Instruction Multiple Data)** operations for parallel data processing.
- 4. **Fault Handling**
 - Supports **HardFault, MemManage, BusFault, and UsageFault** exceptions.
 - Fault escalation to HardFault for unrecoverable errors.
- 5. **Memory Model**
 - **Bit-banding** for atomic bit manipulation in specific memory regions.
 - **Little-endian or big-endian** support (implementation-defined).
- 6. **Synchronization Primitives**
 - **Load-Exclusive (LDREX)** and **Store-Exclusive (STREX)** for atomic operations.
 - Useful for semaphores and multi-core systems.
- 7. **CMSIS (Cortex Microcontroller Software Interface Standard)**
 - Provides standardized APIs for accessing core peripherals and features.
 - Simplifies software development across different Cortex-M4 devices.

Summary

The Cortex-M4 is optimized for embedded applications requiring high performance, low power, and real-time capabilities. Its key strengths include:

- **Efficient processing** with Thumb-2 and optional FPU.
- **Advanced interrupt handling** via NVIC.
- **Flexible memory management** with MPU and bit-banding.
- **DSP extensions** for signal processing tasks.
- **Robust debug and trace** for development.

1. What is ARM Cortex-M?

ARM Cortex-M is a family of 32-bit RISC (Reduced Instruction Set Computer) processor cores designed by ARM for **low-power, cost-sensitive embedded applications**. It is commonly used in microcontrollers for automotive, industrial, consumer, and IoT devices. Key features include:

- Deterministic interrupt handling (great for real-time systems).
- Low-power consumption.
- Thumb-2 instruction set for high code density.

2. Difference between Cortex-M3 and Cortex-M4

Feature	Cortex-M3	Cortex-M4
DSP Instructions	✗ Not available	✓ Available (optimized for DSP tasks)
FPU (Floating Point)	✗ Not available	✓ Optional single-precision FPU
Use Case	General-purpose embedded apps	Signal processing + general applications
Performance	Good	Better for DSP-heavy workloads

3. What is Thumb-2 instruction set?

Thumb-2 is a **mixed 16-bit and 32-bit instruction set** that provides the performance of ARM's 32-bit instructions but with the code density of 16-bit instructions. It allows:

- Smaller binary size (memory savings).
- Efficient performance in embedded systems.
- All Cortex-M processors use only the Thumb-2 instruction set (no full ARM instruction set).

4. What is Harvard vs. Von Neumann architecture in ARM?

- **Harvard Architecture:** Separate memory and buses for **instructions** and **data**. Allows simultaneous access, improving performance.
- **Von Neumann Architecture:** Single memory and bus for both instructions and data. Simpler but can create bottlenecks.

ARM Cortex-M uses a **modified Harvard architecture**:

- Fetch instructions and data separately.
- Unified memory map (single address space) for simplicity.

5. Explain the ARM Cortex-M pipeline stages

Cortex-M processors typically use a **3-stage pipeline**:

1. **Fetch** – Fetch instruction from memory.
2. **Decode** – Decode the instruction to determine the operation and operands.
3. **Execute** – Perform the operation (e.g., arithmetic, memory access).

This pipeline improves performance by allowing overlapping execution of instructions.

6. What is the role of the NVIC in Cortex-M?

NVIC (Nested Vectored Interrupt Controller):

- Manages interrupts in the Cortex-M processors.
- Supports **nested interrupts** (higher priority can preempt lower).
- Integrated into the core for **low-latency interrupt handling**.
- Allows **dynamic priority levels**, **interrupt masking**, and **software-triggered interrupts**.

7. What is the maximum clock speed of Cortex-M3/M4?

- **Cortex-M3:** Up to **100 MHz**, depending on the implementation (some vendors go higher).
- **Cortex-M4:** Typically **100 – 168 MHz**, some variants (e.g., STM32F4) go up to **180 MHz** or more.

Actual max frequency depends on **silicon vendor** and **process technology**.

8. What are the different operating modes in Cortex-M?

Cortex-M has **two main operating modes**:

- **Thread mode:**
 - Executes application code.
 - Can be **privileged** or **unprivileged**.
- **Handler mode:**
 - Used for **exception handling** (interrupts, faults).
 - Always runs in **privileged** mode.

9. What is privileged vs. unprivileged execution mode?

- **Privileged Mode:**
 - Full access to all system resources.
 - Can configure system registers, access protected memory, etc.
- **Unprivileged Mode:**
 - Restricted access.
 - Cannot modify system-critical registers.
 - Used for running user applications to improve system security and stability.

Switching between modes can be controlled via software (e.g., using the CONTROL register or system calls like SVC).

10. What is the MPU (Memory Protection Unit) in Cortex-M?

MPU (Memory Protection Unit):

- Allows **hardware-based memory access control**.
- Can define memory regions with attributes (e.g., read-only, no-execute).

- Helps:
 - Enforce memory boundaries.
 - Prevent faulty/malicious code from corrupting memory.
 - Enable secure embedded applications.
- Available in **Cortex-M3, M4**, and higher models.

1. List the core registers in Cortex-M3/M4

Cortex-M3/M4 processors implement a **set of 16 general-purpose and special-purpose registers** accessible in both thread and handler modes:

Register Name	Description
R0–R12	General-purpose registers
R13	Stack Pointer (SP) — split into MSP and PSP
R14	Link Register (LR) — stores return address
R15	Program Counter (PC) — holds current instruction address
xPSR	Program Status Register: contains APSR, IPSR, and EPSR

Additional control-related registers:

- **CONTROL register** — configures SP (MSP/PSP) and privilege level.
- **PRIMASK, FAULTMASK, BASEPRI** — control interrupt masking levels.

✓ 2. What is the Program Counter (PC) and Link Register (LR)?

- **Program Counter (PC/R15):**
 - Holds the **memory address** of the next instruction to execute.
 - Automatically updated after each instruction fetch.
 - During branching or exception return, it gets explicitly set.
- **Link Register (LR/R14):**
 - Stores the **return address** during a subroutine call or exception.
 - When using BL (Branch with Link) or entering an exception, LR is automatically set.
 - In exceptions, it can also hold special EXC_RETURN values to indicate how to return.

✓ 3. What is the Stack Pointer (SP)? Difference between MSP and PSP

- **Stack Pointer (SP/R13)** points to the **top of the stack** in memory. It is used to:

- Store return addresses, local variables, and register values during function calls or interrupts.

There are **two types of SPs** in Cortex-M:

Stack Pointer	Description
MSP (Main Stack Pointer)	Default stack after reset; used in handler mode (interrupts)
PSP (Process Stack Pointer)	Used in thread mode , typically by application code for separation and security

🧠 **Why two?** This separation allows keeping the OS/privileged stack (MSP) secure, while user applications run on a different, isolated stack (PSP).

✓ 4. What is the Application Program Status Register (APSR)?

- **APSR** contains **condition flags** resulting from arithmetic/logical operations:

Flag	Meaning
N	Negative result
Z	Zero result
C	Carry
V	Overflow
Q	Saturation (DSP overflow, in M4 only)

These flags can influence **conditional instructions** and program flow.

✓ 5. What is the Execution Program Status Register (EPSR)?

- **EPSR** contains **execution state information**, such as:

Field	Purpose
T	Indicates Thumb state (always 1 in Cortex-M)
IT bits	Hold status of IT (If-Then) blocks (conditional instructions)
Reserved	Internal use

Typically, EPSR is not directly written to by software.

✓ 6. What is the Interrupt Program Status Register (IPSR)?

- **IPSR** shows the **exception number** currently being handled.

IPSR Value	Meaning
0	Thread mode (no exception)
1-15	System exceptions (e.g., NMI, HardFault)
16+	External interrupts (IRQ0 and beyond)

It's useful for **debugging** or runtime diagnostics.

✓ 7. What is banked stacking in Cortex-M?

When an **exception occurs**, Cortex-M performs **automatic context saving** (banked stacking):

- Saves registers **R0–R3, R12, LR, PC, and xPSR** onto the **current stack**.
- The processor automatically chooses between **MSP or PSP**, depending on the current mode and configuration.

🧠 This is "banked" because MSP and PSP are **two separate register banks** for stack use, depending on the current operating mode.

✓ 8. How does exception handling work in Cortex-M?

Here's the step-by-step flow:

1. **Exception triggers** (interrupt, fault, etc.).
2. **Processor switches to handler mode** and sets **MSP** (if not already using it).

3. Cortex-M **automatically pushes** a context frame (R0–R3, R12, LR, PC, xPSR).
4. Jumps to the **exception handler** using the vector table.
5. After the ISR runs, the **context is popped** off the stack to resume the pre-interrupt state.

➔ This makes exception entry and return **predictable and efficient**.

✓ 9. What is tail-chaining in interrupts?

- **Tail-chaining** is an optimization when **multiple pending interrupts** occur back-to-back.
- When one ISR completes and another one is pending:
 - Cortex-M **skips the context restore and save** steps between ISRs.
 - **Directly jumps** to the next ISR.

✓ This reduces **latency** and improves **performance**.

✓ 10. What is late-arriving interrupt handling?

A **late-arriving interrupt** is when a **higher-priority interrupt** arrives **while another interrupt is still entering**.

- Cortex-M handles this by checking for **new higher-priority interrupts before completing the stack frame save**.
- If one is found, it **aborts the current ISR entry** and immediately **starts handling the higher-priority one**.

This allows **fast switching to the most urgent task**, maintaining **real-time responsiveness**.

1. What is the difference between IRQ and Exception?

Feature	IRQ (Interrupt Request)	Exception
Source	External hardware peripherals	Internal system events (e.g., faults, system calls)
Examples	GPIO interrupt, USART RX, Timer	HardFault, NMI, SVC, SysTick
Numbering	IRQs start from IRQ0 (exception number 16)	Exceptions have reserved numbers 1–15
Managed by	NVIC (Nested Vectored Interrupt Controller)	System control logic + NVIC

➔ All IRQs are exceptions, but **not all exceptions are IRQs**.

✓ 2. How many interrupt priorities are there in Cortex-M?

- Cortex-M supports **programmable priority levels** using **NVIC**.
- Number of priority levels depends on implementation (usually 8 to 256 levels):
 - Typically **3 to 8 bits** are used for priority (`__NVIC_PRIO_BITS`).
 - For example, if 3 bits → **8 priority levels** (0 = highest, 7 = lowest).

➔ Priorities are **grouped** and can include **preemption** and **subpriority**.

✓ 3. What is priority grouping in NVIC?

- NVIC allows you to split interrupt priority into:
 - **Preempt priority**: Determines if an interrupt can preempt another.
 - **Subpriority**: Resolves priority if two interrupts have the same preempt level.

You configure this using:

```
NVIC_SetPriorityGrouping(NVIC_PRIORITYGROUP_x);
```

Where x defines the split between preemptive and subpriority bits.

➔ Useful in real-time systems to manage nested interrupt behavior.

✓ 4. How to configure an interrupt in Cortex-M?

Steps to configure a peripheral interrupt:

1. **Enable the peripheral interrupt** in the NVIC:

```
NVIC_EnableIRQ(IRQn_Type irq);
```

1. **Set the priority:**

```
NVIC_SetPriority(IRQn_Type irq, uint32_t priority);
```

1. **Configure the peripheral interrupt enable/trigger** (depends on the peripheral, e.g., timer, USART).
2. **Write the ISR handler:**

```
void Your_IRQ_Handler(void){
```

```
    // Clear interrupt flag
```

```
    // Handle logic
```

```
}
```

➔ All interrupt vectors must be mapped in the **vector table**.

✓ 5. What is PendSV and why is it used in RTOS?

- **PendSV (Pending Supervisor Call)** is a **software-triggered interrupt** with **low priority**.
- Designed for **context switching** in **RTOS environments** like FreeRTOS or RTX.

Why use PendSV?

- RTOS tasks are switched only **after** higher-priority ISRs complete.
- PendSV is **pending manually** and executes **after all other interrupts**, ensuring a safe context switch.

SCB->ICSR |= SCB_ICSR_PENDSVSET_Msk; // Trigger PendSV

✓ 6. What is SVC (Supervisor Call)?

- **SVC (Exception #11)** is a software-generated exception used to:
 - Request **privileged operations** from unprivileged code.
 - Call **OS/system services** in RTOS.

Generated using:

assembly

CopyEdit

SVC #<imm8>

➔ On SVC, the processor enters handler mode and executes the **SVC_Handler** function.

✓ 7. What is the SysTick timer and its use?

- **SysTick** is a **24-bit countdown timer** built into all Cortex-M processors.
- Uses:
 - **Periodic interrupts** (e.g., every 1 ms) for timekeeping.
 - **RTOS tick generation** for task scheduling.
 - Delay and timeout functions.

Key registers:

- SYST_CSR: Control
- SYST_RVR: Reload value
- SYST_CVR: Current value

Config example:

SysTick_Config(SystemCoreClock / 1000); // 1ms tick

✓ 8. How does nested interrupt handling work?

- Cortex-M allows **higher-priority interrupts** to **preempt lower-priority ISRs**.
- The NVIC handles:
 - Saving the current context (automatically).

- Switching to the higher-priority ISR.
- Returning via **automatic stacking/unstacking**.

➔ You control nesting behavior via **priority levels** and **priority grouping**.

✓ 9. What is interrupt latency in Cortex-M? How to minimize it?

- **Interrupt latency**: Time from **interrupt assertion** to **start of ISR execution**.
- Typical latency for Cortex-M3/M4: **12 cycles** (best case, no stacking delays).

Ways to minimize latency:

- Keep ISRs **short and efficient**.
- Avoid disabling interrupts for too long (`__disable_irq()`).
- Use **priority levels** effectively.
- **Avoid long critical sections** or heavy processing inside ISRs.
- Enable **tail-chaining** and **late-arrival handling** features (handled automatically by Cortex-M).

✓ 10. What is the Wake-up Interrupt Controller (WIC)?

- **WIC** is a **low-power interrupt wakeup controller** in Cortex-M.
- When the CPU is in **deep sleep** (e.g., WFI or WFE), the main processor logic is off, but **WIC remains active**.
- It detects **wake-up capable interrupts** and **triggers CPU wake-up**.

Key benefits:

- Reduces power usage while still responding to important events.
- Essential for **ultra-low-power embedded systems**.

1. What is the memory map of Cortex-M3/M4?

Cortex-M3/M4 processors use a **linear 4 GB address space** (32-bit) with memory regions defined as follows:

Address Range	Region Name	Description
0x0000_0000–0x1FFF_FFFF	Code (Flash)	Main code execution area (Flash memory)
0x2000_0000–0x3FFF_FFFF	SRAM	On-chip SRAM (data memory)
0x4000_0000–0x5FFF_FFFF	Peripheral	Peripheral registers (GPIO, UART, etc.)
0x6000_0000–0x9FFF_FFFF	External RAM	Off-chip memory (FSMC/FMC)
0xA000_0000–0xDFFF_FFFF	External Devices	External devices
0xE000_0000–0xE00F_FFFF	System Control Space	NVIC, SysTick, SCB, etc.
0xE010_0000–0xFFFF_FFFF	Vendor-specific	Implementation-defined (e.g., debug units)

✓ 2. What are Bit-Banding regions? How do they work?

Bit-Banding allows **atomic bit-level access** to a word in memory using a **single machine instruction**.

Cortex-M defines two bit-band regions:

1. **SRAM Bit-band:** 0x20000000 – 0x200FFFFFFF
2. **Peripheral Bit-band:** 0x40000000 – 0x400FFFFFFF

Each bit in these regions is **aliased** to a **32-bit word** in the bit-band alias region:

- SRAM alias: 0x22000000 – 0x23FFFFFFF
- Peripheral alias: 0x42000000 – 0x43FFFFFFF

How it works:

To access bit *n* at byte *b*:

$\text{alias_address} = \text{alias_base} + (\text{byte_offset} \times 32) + (\text{bit_number} \times 4)$

- ✓ Writing 1 or 0 to the alias address sets/clears the bit.

Use cases:

- Efficient **bit manipulation**
- Atomic **flag updates**
- **Low-power GPIO toggling**

✓ 3. What is Little-Endian vs. Big-Endian in ARM?

- **Little-Endian** (default for Cortex-M):
 - Least significant byte is stored at **lowest address**.
 - Example:
 - 0x12345678 stored as → [0x78][0x56][0x34][0x12]
- **Big-Endian**:
 - Most significant byte is at lowest address.
 - Example:
 - css
 - CopyEdit
 - 0x12345678 stored as → [0x12][0x34][0x56][0x78]

Cortex-M supports **Little-Endian only**, though some ARM cores allow dynamic switching.

✓ 4. How does Flash memory work in Cortex-M?

- **Flash memory** stores **code and constants**.
- Non-volatile: retains data after power-off.
- Usually mapped to 0x00000000 in memory.
- Flash is **slower than SRAM**, so **prefetch buffers** and **instruction cache** help with speed.
- Flash must be **erased before write** (usually by blocks/pages).

In embedded systems, updating firmware or configurations involves:

- **In-application programming (IAP)**
- **Bootloaders**

✓ 5. What is SRAM and how is it used?

- **Static RAM (SRAM)** is **volatile memory** used for:
 - Stack and heap
 - Global/static variables
 - Buffering and data caching

Benefits:

- **Fast access time**
- **Low latency**

Mapped at: 0x20000000 and used in:

- **Run-time data storage**
- Real-time data processing
- DMA destination/source

✓ 6. What is DMA (Direct Memory Access) in Cortex-M?

- **DMA controller** allows data transfer between peripherals and memory **without CPU involvement**.

Advantages:

- Frees CPU for other tasks
- Higher throughput
- Ideal for **UART, ADC, SPI, I2C, etc.**

Key features:

- Memory-to-peripheral or memory-to-memory transfer
- Circular or burst mode
- Interrupts on transfer complete

DMA configuration includes:

- Source/Destination address
- Number of data items
- Control flags (increment, size, priority)

✓ 7. How does Memory Protection (MPU) work?

- **MPU (Memory Protection Unit)** allows:
 - Isolating memory regions
 - Setting access permissions (Read/Write/Execute)
 - Preventing stack overflows, buggy access

MPU features:

- Up to **8 memory regions**
- Define:
 - Base address
 - Size (e.g., 1KB, 4KB)
 - Permissions
 - Cacheability & shareability

Used in:

- **RTOS** for task isolation
- Secure firmware update
- Safety-critical applications

✓ 8. What is the TCM (Tightly Coupled Memory)?

- **TCM** is **ultra-fast memory** closely connected to the CPU core.
- Used in **Cortex-M7**, not typically in M3/M4.
- Provides:
 - **Zero-wait state** access
 - Ideal for **real-time critical code** or **buffers**

Types:

- **ITCM**: Instruction TCM
- **DTCM**: Data TCM

Mapped to: 0x00000000 or 0x20000000 depending on usage.

✓ 9. How to configure GPIO in Cortex-M?

GPIOs are configured by writing to **memory-mapped registers** in the peripheral region (0x4000_0000+).

Steps:

1. **Enable GPIO port clock:**

```
RCC->AHB1ENR |= RCC_AHB1ENR_GPIOAEN;
```

1. **Set pin mode (input/output/alt/analog):**

```
GPIOA->MODER |= (1 << (2*pin)); // Output mode
```

1. **Set output type (push-pull/open-drain):**

```
GPIOA->OTYPER &= ~(1 << pin); // Push-pull
```

1. **Set speed and pull-up/down:**

c

CopyEdit

```
GPIOA->OSPEEDR |= (1 << (2*pin));
```

```
GPIOA->PUPDR &= ~(3 << (2*pin)); // No pull
```

1. **Write to pin:**

```
GPIOA->ODR |= (1 << pin); // Set pin high
```

✓ Abstracted using HAL/LL libraries in real-world applications.

✓ 10. What is the role of the AHB/APB bus?

Cortex-M uses the **AMBA bus** architecture with two main buses:

Bus	Full Form	Use Case
AHB	Advanced High-performance Bus	High-speed peripherals (DMA, Flash, SRAM)
APB	Advanced Peripheral Bus	Low-speed peripherals (UART, GPIO, I2C, SPI)

Key Differences:

- **AHB** supports **burst transfers**, pipelining, and high throughput.
- **APB** is **simpler** and **lower power**, often used for control/status registers.

In microcontrollers (e.g., STM32):

- AHB is the backbone bus.
- APB1 and APB2 connect slower peripherals.

1. What debugging interfaces are supported in Cortex-M? (JTAG, SWD)

Cortex-M supports two main hardware debugging interfaces:

Interface	Description
JTAG (Joint Test Action Group)	4–5 pin traditional debugging interface supporting boundary scan and full debugging features.
SWD (Serial Wire Debug)	2-pin alternative to JTAG (SWDIO, SWCLK), ideal for space-constrained systems with same debugging capabilities.

➡ Most modern Cortex-M MCUs support **SWD** as the default because it's **smaller and faster** than JTAG.

✓ 2. What is CoreSight technology?

ARM CoreSight is a suite of debug and trace components integrated into ARM processors, including Cortex-M. Its goal is to enable:

- **Non-intrusive debugging**
- **Instruction and data tracing**
- **Profiling and code coverage**

Main components include:

- **DAP (Debug Access Port)**
- **SWD/JTAG**
- **ITM (Instrumentation Trace Macrocell)**
- **ETM (Embedded Trace Macrocell)**
- **DWT (Data Watchpoint and Trace)**
- **TPIU (Trace Port Interface Unit)**

➔ **CoreSight** enables deep visibility into embedded applications without significantly affecting performance.

✓ 3. What is ITM (Instrumentation Trace Macrocell)?

- ITM allows **low-latency debug trace messages** from the firmware.
- Used for:
 - printf-style debugging via **SWO**
 - Event logging
 - RTOS-aware debugging

Data sent via **Stimulus Registers**, then output through the **TPIU** over SWO or trace port.

Key advantage:

- Much faster than semihosting or regular UART debugging.

✓ 4. What is ETM (Embedded Trace Macrocell)?

- **ETM** provides **real-time instruction-level tracing**.
- Allows you to reconstruct the **exact instruction flow** the processor executed.
- Non-intrusive and useful for:
 - Performance profiling
 - Trace-based code coverage
 - In-depth debugging

➔ Requires **dedicated trace pins** (TRACEDATA, TRACECLK) and an external debug tool (e.g., Segger J-Trace, Keil ULINKpro).

✓ 5. How does SWO (Serial Wire Output) work?

- SWO is a **single-pin output line** for trace/debug data (part of SWD).
- Used to transmit:

- **ITM messages**
- **Program counter samples**
- **Timestamp info**

🔧 To use SWO:

- Enable ITM, TPIU, and DWT.
- Configure baud rate.
- Output can be read using tools like **ST-Link Utility**, **Ozone**, **Keil µVision**, etc.

✓ 6. What is breakpoint and watchpoint in debugging?

Type	Description
Breakpoint	Halts program execution at a specific code address .
Watchpoint	Triggers on data access (read/write) to a memory location.

Cortex-M has:

- Up to **6 hardware breakpoints** (via FPB - Flash Patch and Breakpoint unit).
- **4 watchpoints** via the **DWT (Data Watchpoint and Trace) unit**.

Used for:

- **Non-intrusive monitoring**
- **Debugging in Flash** where software breakpoints aren't possible.

✓ 7. What is semihosting?

- **Semihosting** is a way for a target application (running on MCU) to communicate with the host debugger.
- Used for:
 - File I/O
 - printf() output
 - Keyboard input

It uses **SVC (Supervisor Call)** instructions to trap to the debugger.

● Downsides:

- **Very slow**
- Breaks **real-time behavior**

- Needs debugger connection (will crash otherwise)

✓ 8. How to use printf() in embedded ARM?

Several options:

Method	Description
Semihosting	Use printf and redirect via SVC calls. Debugger must be connected.
ITM/SWO	Fast, efficient debug prints over SWO pin using ITM_SendChar().
UART	Traditional method: redirect printf() to UART with putchar() override.

Example using SWO:

```
ITM_SendChar('A'); // Sends character over SWO
```

For UART:

```
int __io_putchar(int ch) {
    HAL_UART_Transmit(&huart2, (uint8_t*)&ch, 1, HAL_MAX_DELAY);
    return ch;
}
```

✓ 9. What is HardFault? How to debug it?

A **HardFault** occurs when the system encounters a critical failure it can't recover from.

Examples:

- Invalid memory access
- Stack overflow
- Undefined instruction
- Division by zero (if trap enabled)

💡 **Debugging steps:**

- Enable **fault handlers** (MemManage, BusFault, UsageFault)
- Read **stack frame** on fault

- Use the CFSR, HFSR, BFAR, MMFAR registers for fault details
- Use a **HardFault handler** to log values:

```
void HardFault_Handler(void) {

    __asm("BKPT #0"); // Triggers debugger breakpoint

}
```

✓ 10. What is BusFault, MemManage Fault, UsageFault?

These are **precise fault types** to help identify bugs:

Fault Type	Trigger Conditions
BusFault	Memory access errors (e.g., invalid peripherals, bus errors)
MemManage Fault	Access to restricted memory via MPU, or stack errors
UsageFault	Illegal instructions, undefined behavior (e.g., divide by zero, unaligned access)

To enable handlers:

```
SCB->SHCSR |= SCB_SHCSR_USGFAULTENA_Msk |

             SCB_SHCSR_BUSFAULTENA_Msk |

             SCB_SHCSR_MEMFAULTENA_Msk;
```

➔ These faults are more **granular than HardFault** and should always be used in **development and testing**.

✓ 1. How does an RTOS work on Cortex-M?

An **RTOS (Real-Time Operating System)** manages tasks, scheduling, and synchronization in time-sensitive embedded systems.

On **Cortex-M**, an RTOS utilizes:

- **SysTick Timer** (for periodic tick interrupts)
- **PendSV** (for context switching)
- **SVC** (for system calls)

- **NVIC** (for interrupt control)

The RTOS divides execution into **threads/tasks**, managing when each runs based on priorities and states (Ready, Running, Blocked, Suspended).

Popular RTOSs: **FreeRTOS, CMSIS-RTOS, Zephyr, RTX.**

✓ 2. What is context switching in ARM?

Context switching is the process of saving the current task's context (CPU registers, stack pointer, program counter, etc.) and restoring another task's context.

On Cortex-M:

- **Hardware saves part of context** (R0–R3, R12, LR, PC, xPSR) automatically during exception entry.
- RTOS uses **PendSV** to complete saving and restoring full context (R4–R11, etc.).

This allows smooth switching between tasks without interfering with each other's execution.

✓ 3. Why is PendSV used in RTOS scheduling?

PendSV (Pending Supervisor Call) is a low-priority, software-triggered interrupt.

Why RTOS uses it:

- It is **deliberately lowest priority**, so it won't interrupt any other IRQ.
- It's perfect for **deferring context switch** after a higher-priority ISR finishes.
- RTOS triggers it when a task needs to yield, or a higher-priority task becomes ready.

🔄 **SysTick** triggers the time slice → **Scheduler decides next task** → **PendSV is triggered** → **Context switch happens.**

✓ 4. What is thread mode vs. handler mode?

Mode	Description
Thread Mode	Executes normal program tasks (main(), threads). It's where the OS tasks run.
Handler Mode	Executes exception handlers (ISRs, SysTick, SVC, PendSV, etc.)

- Cortex-M always **starts in thread mode**.
- On interrupts or faults, it **switches to handler mode**.

- Some instructions (like SVC) explicitly enter handler mode.

✓ 5. How does preemptive scheduling work in Cortex-M?

Preemptive scheduling allows higher-priority tasks to interrupt lower-priority ones.

In Cortex-M:

1. RTOS sets up **SysTick** to generate periodic ticks (e.g., every 1 ms).
2. On SysTick interrupt, RTOS checks if a higher-priority task is ready.
3. If yes, it triggers **PendSV** to do a context switch.

➔ The **NVIC and exception priorities** allow this to happen **efficiently and safely**, with **deterministic timing**.

✓ 6. What is stack overflow protection in RTOS?

Each task in an RTOS gets its **own stack**.

Stack overflow happens when a task exceeds its allocated stack memory, potentially corrupting other memory regions.

Solutions:

- Use **MPU (Memory Protection Unit)** to guard stack boundaries.
- Enable **RTOS stack checking** (e.g., in FreeRTOS: configCHECK_FOR_STACK_OVERFLOW).
- Fill stack with a known pattern (like 0xA5) and check usage periodically.

✓ 7. How to implement mutex/semaphore in Cortex-M?

These are synchronization primitives in RTOS to manage resource access.

- **Mutex:** Prevents simultaneous access to shared resources (with ownership and priority inheritance).
- **Semaphore:** Signals between tasks or ISRs.

In FreeRTOS:

```
SemaphoreHandle_t xSemaphore = xSemaphoreCreateBinary();
```

```
xSemaphoreTake(xSemaphore, portMAX_DELAY); // Wait
```

```
xSemaphoreGive(xSemaphore); // Release
```

In CMSIS-RTOS2:

```
osMutexAcquire(mutex_id, osWaitForever);
```

```
osMutexRelease(mutex_id);
```

⚠ Always **disable interrupts or use atomic operations** when accessing critical sections without RTOS.

✓ 8. What is priority inversion and how to avoid it?

Priority Inversion occurs when:

- A **low-priority task holds a resource**
- A **high-priority task is blocked waiting for it**
- A **medium-priority task** runs and delays the release

Example:

cpp

CopyEdit

High-Priority → waiting on mutex

↓

Low-Priority → owns mutex

↓

Medium-Priority → runs indefinitely, delaying Low → blocks High!

Solution:

Use **Priority Inheritance**:

- When a low-priority task holds a mutex, and a higher-priority task waits, the RTOS **temporarily boosts** the priority of the low-priority task.

FreeRTOS and CMSIS-RTOS support this.

✓ 9. What is tickless mode in RTOS?

In **tickless mode**, the periodic SysTick timer is **disabled during idle** to reduce power consumption.

How it works:

- Instead of firing every 1ms, the RTOS **sleeps longer** and wakes up only when needed (next task delay/timeout/interrupt).
- Great for **battery-powered devices**.

Enabling in FreeRTOS:

```
#define configUSE_TICKLESS_IDLE 1
```

Needs careful timer configuration and accurate wake-up from low-power states.

✓ 10. How to measure task execution time in Cortex-M?

Several ways:

A. DWT Cycle Counter (ARM Debug unit):

```
DWT->CTRL |= DWT_CTRL_CYCCNTENA_Msk; // Enable cycle counter
```

```
DWT->CYCCNT = 0; // Reset
```

```
uint32_t start = DWT->CYCCNT;
```

```
// Task code here
```

```
uint32_t end = DWT->CYCCNT;
```

```
printf("Cycles: %lu\n", end - start);
```

B. GPIO toggling + Oscilloscope

- Toggle a pin before/after task
- Measure on oscilloscope or logic analyzer

C. RTOS Trace Tools


- Tools like **Segger SystemView**, **FreeRTOS+Trace**, or **Keil Event Recorder** can track execution time per task.

1. What are the low-power modes in Cortex-M?

Cortex-M processors (M0/M3/M4/M7) support various low-power modes to **reduce power consumption** by controlling CPU, memory, and peripheral activity.

The core power modes are:

Mode	Description
Sleep Mode	CPU is halted, peripherals can run. Wakes up on interrupt.
Deep Sleep Mode	CPU + more system components halted, clocks stopped. Lower power.
Standby / Shutdown / Stop Mode (<i>SoC-specific</i>)	Deepest sleep. SRAM or even register content lost. Needs full reboot or partial recovery.

 The exact naming and behavior (*Stop, Standby, Shutdown*) vary based on the MCU vendor (like *ST, NXP, etc.*), but they are built on top of ARM's *Sleep* and *Deep Sleep*.

✓ 2. What is Sleep, Deep Sleep, and Standby mode?

Mode	Power Saving	What is OFF	Wake-up Sources
Sleep	Moderate	CPU halted, clocks running	Any interrupt
Deep Sleep	High	CPU halted, system clocks off	External IRQs, RTC, timers
Standby (<i>MCU-dependent</i>)	Maximum	Core + SRAM off (can lose context)	Wake-up pin, RTC, reset

 Key control:

- **SLEEPDEEP bit** in System Control Register (SCR):
 - 0: Sleep
 - 1: Deep Sleep
- Enter using **WFI** or **WFE** instruction.

✓ 3. How to wake up from low-power mode?

- ✓ Wake-up sources depend on MCU and mode:

- **Interrupts (NVIC):** Timer, GPIO, USART, RTC, etc.
- **Reset:** Manual reset, watchdog
- **Events:** External pins, peripherals

For example, to enable a GPIO interrupt to wake from deep sleep:

```
NVIC_EnableIRQ(EXTIO_IRQn); // Enable external interrupt
```

```
__WFI();           // Enter sleep
```

Some MCUs have a dedicated **Wake-Up Controller (WIC)** to handle wake-up during deep sleep.

✓ 4. What is WFI (Wait for Interrupt) and WFE (Wait for Event)?

These are **ARM Cortex-M instructions** that transition the processor into a low-power mode.

Instruction	Behavior
WFI (Wait for Interrupt)	Puts the core to sleep until any interrupt occurs.
WFE (Wait for Event)	Waits for event signal (more flexible, works with SEV). Can be used in multi-core scenarios or low-power sync.

⚠ WFI is more commonly used in bare-metal and RTOS systems.

Example (Sleep mode entry):

```
SCB->SCR &= ~SCB_SCR_SLEEPDEEP_Msk; // Sleep mode
```

```
__WFI();           // Enter sleep
```

Example (Deep Sleep mode entry):

```
SCB->SCR |= SCB_SCR_SLEEPDEEP_Msk; // Deep Sleep
```

```
__WFI();           // Enter deep sleep
```

✓ 5. How does clock gating work in Cortex-M?

Clock gating is a technique where clocks to unused modules (CPU, timers, peripherals) are **disabled to save power**.

Key aspects:

- Managed by the **MCU's power control and RCC/PMC unit** (not by ARM core itself).
- Typically MCU vendors (like ST, NXP) provide APIs or registers to:
 - Enable/disable peripheral clocks
 - Put specific modules into low-power mode
 - Reduce or stop system clocks

Example (STM32 HAL):

```
__HAL_RCC_GPIOA_CLK_DISABLE(); // Gate clock to GPIOA
```



By gating the clock, dynamic power consumption is reduced because switching activity in logic stops.

Summary of Low-Power Features

Feature	Purpose
Sleep / Deep Sleep	Save CPU + system power during idle
WFI / WFE	Enter sleep modes based on system state
Clock gating	Disable unused hardware clocks
Wake-up via IRQs or Events	Resume execution on need
MPU + SleepOnExit	Optimize idle + secure transitions

1. What is TrustZone in ARMv8-M?

TrustZone is a **security extension** in **ARMv8-M** architecture (used in Cortex-M23 and Cortex-M33) that allows separation of software into two **execution environments**:

-  **Secure world**: Trusted, sensitive code (e.g., cryptography, keys, secure boot)
-  **Non-Secure world**: Normal application code

 Key Features:

- Secure/Non-Secure **Memory partitioning**
- Secure/Non-Secure **peripheral access control**
- Context-switching between the two environments
- **NSC (Non-Secure Callable)** regions to safely expose secure APIs

✓ Benefits:

- Improved security in IoT and embedded systems
- Enables **Trusted Execution Environment (TEE)** on microcontrollers
- Helps with **security certifications** (e.g., PSA Certified)

✓ 2. What is Secure and Non-Secure Mode in ARM?

These are the two **operating states** introduced by TrustZone in ARMv8-M:

Mode	Description
Secure Mode	Has access to all system memory and peripherals
Non-Secure Mode	Restricted access, cannot access secure memory/peripherals

The transition between these modes is managed by:

- **Secure Gateway (SG) instructions**
- **Non-Secure Callable (NSC)** regions
- The **Security Attribution Unit (SAU)** or **Implementation Defined Attribution Unit (IDAU)**

Example Use:

- Secure firmware (crypto libraries, keys) runs in secure mode.
- Main application (UI, data logging) runs in non-secure mode.

✓ 3. How does dual-core (Cortex-M4 + Cortex-M0) work?

Some MCUs (e.g., **NXP LPC55S69**, **Nordic nRF5340**, **STM32WB**) use **dual-core systems**, usually pairing:

- **Cortex-M4**: Application processor (runs the main application)
- **Cortex-M0(+)**: Coprocessor (handles Bluetooth stack, sensor reading, real-time tasks)

💡 Operation Model:

- Both cores are **independent but share memory**.

- Communication happens via:
 - **Shared RAM**
 - **Inter-Processor Communication (IPC)** / Mailbox
 - **Hardware semaphores**

Common Roles:

Core	Typical Role
Cortex-M4	App logic, UI, control
Cortex-M0+	BLE stack, sensor polling, secure tasks

These systems are commonly used in **wireless** or **secure edge devices**.

✓ 4. What is cache memory in ARM?

Cache is a small, fast memory used to store recently accessed **instructions or data** to speed up access times.

While basic Cortex-M3/M4 cores **don't include a cache**, higher-end Cortex-M7 and newer ARMv8-M cores like Cortex-M33 **do support caches**.

Cache Types:

- **I-Cache (Instruction Cache)**: Stores recently fetched instructions
- **D-Cache (Data Cache)**: Stores recently accessed data

Cache Policies:

- Write-through
- Write-back
- Cacheable / Non-cacheable regions

Why is cache important?

- Boosts performance by reducing wait states for Flash access
- Helps run high-throughput algorithms efficiently

⚠ Must manage cache consistency with DMA, peripherals, etc. (e.g., invalidate, clean operations)

✓ 5. How to implement bootloader in Cortex-M?

A **bootloader** is a small firmware program that runs before the main application and is responsible for:

- **Firmware upgrade**
- **Integrity verification**
- **Secure boot** (optional)

Steps to Implement a Bootloader:

1. Partition Flash Memory

plaintext

CopyEdit

```
+-----+ 0x08000000
```

```
| Bootloader (16 KB) |
```

```
+-----+ 0x08004000
```

```
| Main Application  |
```

```
+-----+
```

2. Startup Configuration

- Bootloader sets **vector table offset** for application using SCB->VTOR:

```
SCB->VTOR = APPLICATION_ADDRESS;
```

- Set MSP (Main Stack Pointer) and jump to reset handler of application:

```
typedef void (*app_entry_t)(void);
```

```
uint32_t app_sp = *(volatile uint32_t*)APPLICATION_ADDRESS;
```

```
uint32_t app_reset = *(volatile uint32_t*)(APPLICATION_ADDRESS + 4);
```

```
__set_MSP(app_sp);
```

```
((app_entry_t)app_reset)();
```

3. Add Flash Programming Capability

- Receive firmware over UART, USB, BLE, etc.
- Erase and write application flash sectors

4. Optional: Add Security

- Validate firmware using **checksum or cryptographic signature**
- Use **TrustZone or MPU** to isolate bootloader from application

✔ Summary Table

Feature	Cortex-M Support
TrustZone	ARMv8-M (M23, M33)
Secure / Non-Secure Modes	Yes (via SAU/IDAU)
Dual-Core (M4 + M0+)	Vendor-specific (ST, NXP)
Cache	Cortex-M7, Cortex-M33+
Bootloader	Software-implemented